

API Documentation

To connect to a InfluxDB, you must create a [InfluxDBClient](#) object. The default configuration connects to InfluxDB on `localhost` with the default ports. The below instantiation statements are all equivalent:

```
from influxdb import InfluxDBClient

# using Http
client = InfluxDBClient(database='dbname')
client = InfluxDBClient(host='127.0.0.1', port=8086, database='dbname')
client = InfluxDBClient(host='127.0.0.1', port=8086, username='root',
password='root', database='dbname')

# using UDP
client = InfluxDBClient(host='127.0.0.1', database='dbname', use_udp=True,
udp_port=4444)
```

To write pandas DataFrames or to read data into a pandas DataFrame, use a [DataFrameClient](#) object. These clients are initiated in the same way as the [InfluxDBClient](#):

```
from influxdb import DataFrameClient

client = DataFrameClient(host='127.0.0.1', port=8086, username='root',
password='root', database='dbname')
```

Note

Only when using UDP (`use_udp=True`) the connection is established.

InfluxDBClient

```
class influxdb.InfluxDBClient(host=u'localhost', port=8086, username=u'root',
password=u'root', database=None, ssl=False, verify_ssl=False, timeout=None, retries=3,
use_udp=False, udp_port=4444, proxies=None, pool_size=10, path=u'', cert=None, gzip=False,
session=None, headers=None, socket_options=None)
```

InfluxDBClient primary client object to connect InfluxDB.

The [InfluxDBClient](#) object holds information necessary to connect to InfluxDB. Requests can be made to InfluxDB directly through the client.

The client supports the use as a [context manager](#).

- | | |
|-------------------|--|
| Parameters | <ul style="list-style-type: none">• host (<i>str</i>) – hostname to connect to InfluxDB, defaults to ‘localhost’• port (<i>int</i>) – port to connect to InfluxDB, defaults to 8086• username (<i>str</i>) – user to connect, defaults to ‘root’• password (<i>str</i>) – password of the user, defaults to ‘root’• pool_size (<i>int</i>) – urllib3 connection pool size, defaults to 10.• database (<i>str</i>) – database name to connect to, defaults to None• ssl (<i>bool</i>) – use https instead of http to connect to InfluxDB, defaults to False• verify_ssl (<i>bool</i>) – verify SSL certificates for HTTPS requests, defaults |
| : | |

to False

- **timeout** (*int*) – number of seconds Requests will wait for your client to establish a connection, defaults to None
- **retries** (*int*) – number of attempts your client will make before aborting, defaults to 3 0 - try until success 1 - attempt only once (without retry) 2 - maximum two attempts (including one retry) 3 - maximum three attempts (default option)
- **use_udp** (*bool*) – use UDP to connect to InfluxDB, defaults to False
- **udp_port** (*int*) – UDP port to connect to InfluxDB, defaults to 4444
- **proxies** (*dict*) – HTTP(S) proxy to use for Requests, defaults to {}
- **path** (*str*) – path of InfluxDB on the server to connect, defaults to ''
- **cert** (*str*) – Path to client certificate information to use for mutual TLS authentication. You can specify a local cert to use as a single file containing the private key and the certificate, or as a tuple of both files' paths, defaults to None
- **gzip** (*bool*) – use gzip content encoding to compress requests
- **session** (*requests.Session*) – allow for the new client request to use an existing requests Session, defaults to None
- **headers** (*dict*) – headers to add to Requests, will add 'Content-Type' and 'Accept' unless these are already present, defaults to {}
- **socket_options** (*list*) – use custom tcp socket options, If not specified, then defaults are loaded from `HTTPConnection.default_socket_options`

Raises: **ValueError** – if cert is provided but ssl is disabled (set to False)

`alter_retention_policy(name, database=None, duration=None, replication=None, default=None, shard_duration=None)`

Modify an existing retention policy for a database.

- Parameters**
- :
- **name** (*str*) – the name of the retention policy to modify
 - **database** (*str*) – the database for which the retention policy is modified. Defaults to current client's database
 - **duration** (*str*) – the new duration of the existing retention policy. Durations such as 1h, 90m, 12h, 7d, and 4w, are all supported and mean 1 hour, 90 minutes, 12 hours, 7 day, and 4 weeks, respectively. For infinite retention, meaning the data will never be deleted, use 'INF' for duration. The minimum retention period is 1 hour.
 - **replication** (*int*) – the new replication of the existing retention policy
 - **default** (*bool*) – whether or not to set the modified policy as default
 - **shard_duration** (*str*) – the shard duration of the retention policy. Durations such as 1h, 90m, 12h, 7d, and 4w, are all supported and mean 1 hour, 90 minutes, 12 hours, 7 day, and 4 weeks, respectively. Infinite retention is not supported. As a workaround, specify a "1000w" duration to achieve an extremely long shard group duration. The minimum shard group duration is 1 hour.

Note

at least one of duration, replication, or default flag should be set. Otherwise the operation will fail.

`close()`

Close http session.

`create_continuous_query(name, select, database=None, resample_opts=None)`

Create a continuous query for a database.

- Parameters:**
- **name** (*str*) – the name of continuous query to create
 - **select** (*str*) – select statement for the continuous query
 - **database** (*str*) – the database for which the continuous query is created. Defaults to current client's database
 - **resample_opts** (*str*) – resample options

Example:

```
>> select_clause = 'SELECT mean("value") INTO "cpu_mean" ' \
...                 'FROM "cpu" GROUP BY time(1m)'
>> client.create_continuous_query(
...     'cpu_mean', select_clause, 'db_name', 'EVERY 10s FOR 2m'
... )
>> client.get_list_continuous_queries()
[
    {
        'db_name': [
            {
                'name': 'cpu_mean',
                'query': 'CREATE CONTINUOUS QUERY "cpu_mean" '
                        'ON "db_name" '
                        'RESAMPLE EVERY 10s FOR 2m '
                        'BEGIN SELECT mean("value") '
                        'INTO "cpu_mean" FROM "cpu" '
                        'GROUP BY time(1m) END'
            }
        ]
    }
]
```

`create_database(dbname)`

Create a new database in InfluxDB.

Parameters: **dbname** (*str*) – the name of the database to create

`create_retention_policy(name, duration, replication, database=None, default=False, shard_duration=u'0s')`

Create a retention policy for a database.

- Parameters**
- **name** (*str*) – the name of the new retention policy
 - **duration** (*str*) – the duration of the new retention policy. Durations such as 1h, 90m, 12h, 7d, and 4w, are all supported and

mean 1 hour, 90 minutes, 12 hours, 7 day, and 4 weeks, respectively. For infinite retention - meaning the data will never be deleted - use 'INF' for duration. The minimum retention period is 1 hour.

- **replication** (*str*) – the replication of the retention policy
- **database** (*str*) – the database for which the retention policy is created. Defaults to current client's database
- **default** (*bool*) – whether or not to set the policy as default
- **shard_duration** (*str*) – the shard duration of the retention policy. Durations such as 1h, 90m, 12h, 7d, and 4w, are all supported and mean 1 hour, 90 minutes, 12 hours, 7 day, and 4 weeks, respectively. Infinite retention is not supported. As a workaround, specify a "1000w" duration to achieve an extremely long shard group duration. Defaults to "0s", which is interpreted by the database to mean the default value given the duration. The minimum shard group duration is 1 hour.

`create_user(username, password, admin=False)`

Create a new user in InfluxDB.

- Parameters:**
- **username** (*str*) – the new username to create
 - **password** (*str*) – the password for the new user
 - **admin** (*boolean*) – whether the user should have cluster administration privileges or not

`delete_series(database=None, measurement=None, tags=None)`

Delete series from a database.

Series must be filtered by either measurement and tags. This method cannot be used to delete all series, use *drop_database* instead.

- Parameters:**
- **database** (*str*) – the database from which the series should be deleted, defaults to client's current database
 - **measurement** (*str*) – Delete all series from a measurement
 - **tags** (*dict*) – Delete all series that match given tags

`drop_continuous_query(name, database=None)`

Drop an existing continuous query for a database.

- Parameters:**
- **name** (*str*) – the name of continuous query to drop
 - **database** (*str*) – the database for which the continuous query is dropped. Defaults to current client's database

`drop_database(dbname)`

Drop a database from InfluxDB.

- Parameters:** **dbname** (*str*) – the name of the database to drop

`drop_measurement(measurement)`

Drop a measurement from InfluxDB.

Parameters: **measurement** (*str*) – the name of the measurement to drop
`drop_retention_policy(name, database=None)`

Drop an existing retention policy for a database.

Parameters:

- **name** (*str*) – the name of the retention policy to drop
- **database** (*str*) – the database for which the retention policy is dropped. Defaults to current client's database

`drop_user(username)`

Drop a user from InfluxDB.

Parameters: **username** (*str*) – the username to drop
classmethod `from_dsn(dsn, **kwargs)`

Generate an instance of `InfluxDBClient` from given data source name.

Return an instance of [InfluxDBClient](#) from the provided data source name. Supported schemes are “influxdb”, “https+influxdb” and “udp+influxdb”. Parameters for the [InfluxDBClient](#) constructor may also be passed to this method.

Parameters:

- **dsn** (*string*) – data source name
- **kwargs** (*dict*) – additional parameters for *InfluxDBClient*

Raises: **ValueError** – if the provided DSN has any unexpected values

Example:

```
>> cli = InfluxDBClient.from_dsn('influxdb://username:password@\nlocalhost:8086/databasename', timeout=5)\n>> type(cli)\n<class 'influxdb.client.InfluxDBClient'\n>> cli = InfluxDBClient.from_dsn('udp+influxdb://username:pass@\nlocalhost:8086/databasename', timeout=5, udp_port=159)\n>> print('{0._baseurl} - {0.use_udp} {0.udp_port}'.format(cli))\nhttp://localhost:8086 - True 159
```

Note

parameters provided in **kwargs** may override dsn parameters

Note

when using “udp+influxdb” the specified port (if any) will be used for the TCP connection; specify the UDP port with the additional *udp_port* parameter (cf. examples).

`get_list_continuous_queries()`

Get the list of continuous queries in InfluxDB.

Returns: all CQs in InfluxDB

Return type: list of dictionaries

Example:

```
>> cqs = client.get_list_cqs()
>> cqs
[
  {
    u'db1': []
  },
  {
    u'db2': [
      {
        u'name': u'vampire',
        u'query': u'CREATE CONTINUOUS QUERY vampire ON '
                  'mydb BEGIN SELECT count(dracula) INTO '
                  'mydb.autogen.all_of_them FROM '
                  'mydb.autogen.one GROUP BY time(5m) END'
      }
    ]
  }
]
```

`get_list_database()`

Get the list of databases in InfluxDB.

Returns: all databases in InfluxDB

Return type: list of dictionaries

Example:

```
>> dbs = client.get_list_database()
>> dbs
[{u'name': u'db1'}, {u'name': u'db2'}, {u'name': u'db3'}]
```

`get_list_measurements()`

Get the list of measurements in InfluxDB.

Returns: all measurements in InfluxDB

Return type: list of dictionaries

Example:

```
>> dbs = client.get_list_measurements()
>> dbs
[{u'name': u'measurements1'},
 {u'name': u'measurements2'},
 {u'name': u'measurements3'}]
```

`get_list_privileges(username)`

Get the list of all privileges granted to given user.

Parameters: **username** (*str*) – the username to get privileges of

Returns: all privileges granted to given user

Return type: list of dictionaries

Example:

```
>> privileges = client.get_list_privileges('user1')
>> privileges
[{'u'privilege': u'WRITE', u'database': u'db1'},
 {'u'privilege': u'ALL PRIVILEGES', u'database': u'db2'},
 {'u'privilege': u'NO PRIVILEGES', u'database': u'db3'}]
```

`get_list_retention_policies(database=None)`

Get the list of retention policies for a database.

Parameters: **database** (*str*) – the name of the database, defaults to the client's current database

Returns: all retention policies for the database

Return type: list of dictionaries

Example:

```
>> ret_policies = client.get_list_retention_policies('my_db')
>> ret_policies
[{'u'default': True,
  u'duration': u'0',
  u'name': u'default',
  u'replicaN': 1}]
```

`get_list_series(database=None, measurement=None, tags=None)`

Query SHOW SERIES returns the distinct series in your database.

FROM and WHERE clauses are optional.

Parameters:

- **measurement** – Show all series from a measurement
- **tags** – Show all series that match given tags
- **database** (*str*) – the database from which the series should be shows, defaults to client's current database

`get_list_users()`

Get the list of all users in InfluxDB.

Returns: all users in InfluxDB

Return type: list of dictionaries

Example:

```
>> users = client.get_list_users()
>> users
[{'u'admin': True, u'user': u'user1'},
 {'u'admin': False, u'user': u'user2'},
 {'u'admin': False, u'user': u'user3'}]
```

`grant_admin_privileges(username)`

Grant cluster administration privileges to a user.

Parameters: **username** (*str*) – the username to grant privileges to
Note

Only a cluster administrator can create/drop databases and manage users.

`grant_privilege(privilege, database, username)`

Grant a privilege on a database to a user.

Parameters:

- **privilege** (*str*) – the privilege to grant, one of 'read', 'write' or 'all'. The string is case-insensitive
- **database** (*str*) – the database to grant the privilege on
- **username** (*str*) – the username to grant the privilege to

`ping()`

Check connectivity to InfluxDB.

Returns: The version of the InfluxDB the client is connected to
`query(query, params=None, bind_params=None, epoch=None, expected_response_code=200, database=None, raise_errors=True, chunked=False, chunk_size=0, method=u'GET')`

Send a query to InfluxDB.

Danger

In order to avoid injection vulnerabilities (similar to [SQL injection](#) vulnerabilities), do not directly include untrusted data into the `query` parameter, use `bind_params` instead.

Parameters:

- **query** (*str*) – the actual query string
- **params** (*dict*) – additional parameters for the request, defaults to {}
- **bind_params** (*dict*) – bind parameters for the query: any variable in the query written as '\$var_name' will be replaced with `bind_params['var_name']`. Only works in the `WHERE` clause and takes precedence over `params['params']`
- **epoch** (*str*) – response timestamps to be in epoch format either 'h', 'm', 's', 'ms', 'u', or 'ns', defaults to *None* which is RFC3339 UTC format with nanosecond precision
- **expected_response_code** (*int*) – the expected status code of response, defaults to 200
- **database** (*str*) – database to query, defaults to *None*
- **raise_errors** (*bool*) – Whether or not to raise exceptions when InfluxDB returns errors, defaults to *True*
- **chunked** (*bool*) – Enable to use chunked responses from InfluxDB. With `chunked` enabled, one `ResultSet` is returned per chunk containing all results within that chunk
- **chunk_size** (*int*) – Size of each chunk to tell InfluxDB to use.
- **method** (*str*) – the HTTP method for the request, defaults to `GET`

Returns: the queried data

Return [ResultSet](#)

type:

`request(url, method=u'GET', params=None, data=None, stream=False, expected_response_code=200, headers=None)`

Make a HTTP request to the InfluxDB API.

Parameters:

- **url** (*str*) – the path of the HTTP request, e.g. write, query, etc.
- **method** (*str*) – the HTTP method for the request, defaults to GET
- **params** (*dict*) – additional parameters for the request, defaults to None
- **data** (*str*) – the data of the request, defaults to None
- **stream** (*bool*) – True if a query uses chunked responses
- **expected_response_code** (*int*) – the expected response code of the request, defaults to 200
- **headers** (*dict*) – headers to add to the request

Returns: the response from the request

Return `requests.Response`

type:

Raises:

- [InfluxDBServerError](#) – if the response code is any server error code (5xx)
- [InfluxDBClientError](#) – if the response code is not the same as *expected_response_code* and is not a server error code

`revoke_admin_privileges(username)`

Revoke cluster administration privileges from a user.

Parameters: **username** (*str*) – the username to revoke privileges from
Note

Only a cluster administrator can create/ drop databases and manage users.

`revoke_privilege(privilege, database, username)`

Revoke a privilege on a database from a user.

Parameters:

- **privilege** (*str*) – the privilege to revoke, one of 'read', 'write' or 'all'. The string is case-insensitive
- **database** (*str*) – the database to revoke the privilege on
- **username** (*str*) – the username to revoke the privilege from

`send_packet(packet, protocol=u'json', time_precision=None)`

Send an UDP packet.

Parameters:

- **packet** ((if protocol is 'json') *dict* (if protocol is 'line') *list of line*

- protocol strings*) – the packet to be sent
- **protocol** (*str*) – protocol of input data, either ‘json’ or ‘line’
- **time_precision** (*str*) – Either ‘s’, ‘m’, ‘ms’ or ‘u’, defaults to None

`set_user_password(username, password)`

Change the password of an existing user.

- Parameters:**
- **username** (*str*) – the username who’s password is being changed
 - **password** (*str*) – the new password for the user

`switch_database(database)`

Change the client’s database.

Parameters: **database** (*str*) – the name of the database to switch to
`switch_user(username, password)`

Change the client’s username.

- Parameters:**
- **username** (*str*) – the username to switch to
 - **password** (*str*) – the password for the username

`write(data, params=None, expected_response_code=204, protocol=u'json')`

Write data to InfluxDB.

- Parameters:**
- **data** – the data to be written
 - **params** (*dict*) – additional parameters for the request, defaults to None
 - **expected_response_code** (*int*) – the expected response code of the write operation, defaults to 204
 - **protocol** (*str*) – protocol of input data, either ‘json’ or ‘line’

Returns: True, if the write operation is successful

Return type: bool

`write_points(points, time_precision=None, database=None, retention_policy=None, tags=None, batch_size=None, protocol=u'json', consistency=None)`

Write to multiple time series names.

- Parameters:**
- **time_precision** (*str*) – Either ‘s’, ‘m’, ‘ms’ or ‘u’, defaults to None
 - **database** (*str*) – the database to write the points to. Defaults to the client’s current database
 - **tags** (*dict*) – a set of key-value pairs associated with each point. Both keys and values must be strings. These are shared tags and

- will be merged with point-specific tags, defaults to None
- **retention_policy** (*str*) – the retention policy for the points. Defaults to None
- **batch_size** (*int*) – value to write the points in batches instead of all at one time. Useful for when doing data dumps from one database to another or when doing a massive write operation, defaults to None
- **protocol** (*str*) – Protocol for writing data. Either 'line' or 'json'.
- **consistency** (*str*) – Consistency for the points. One of {'any', 'one', 'quorum', 'all'}.

Returns: True, if the operation is successful

Return type: bool

Note

if no retention policy is specified, the default retention policy for the database is used

DataFrameClient

```
class influxdb.DataFrameClient(host=u'localhost', port=8086, username=u'root',
password=u'root', database=None, ssl=False, verify_ssl=False, timeout=None, retries=3,
use_udp=False, udp_port=4444, proxies=None, pool_size=10, path=u'', cert=None, gzip=False,
session=None, headers=None, socket_options=None)
```

DataFrameClient instantiates InfluxDBClient to connect to the backend.

The DataFrameClient object holds information necessary to connect to InfluxDB. Requests can be made to InfluxDB directly through the client. The client reads and writes from pandas DataFrames.

```
EPOCH = Timestamp('1970-01-01 00:00:00+0000', tz='UTC')
query(query, params=None, bind_params=None, epoch=None,
expected_response_code=200, database=None, raise_errors=True, chunked=False,
chunk_size=0, method=u'GET', dropna=True, data_frame_index=None)
```

Query data into a DataFrame.

Danger

In order to avoid injection vulnerabilities (similar to [SQL injection](#) vulnerabilities), do not directly include untrusted data into the query parameter, use bind_params instead.

- Parameters:**
- **bind_params** – bind parameters for the query: any variable in the query written as '\$var_name' will be replaced with bind_params['var_name']. Only works in the WHERE clause and takes precedence over params['params']
 - **epoch** – response timestamps to be in epoch format either 'h', 'm', 's', 'ms', 'u', or 'ns', defaults to None which is RFC3339

UTC format with nanosecond precision

- **expected_response_code** – the expected status code of response, defaults to 200
- **database** – database to query, defaults to None
- **raise_errors** – Whether or not to raise exceptions when InfluxDB returns errors, defaults to True
- **chunked** – Enable to use chunked responses from InfluxDB. With chunked enabled, one ResultSet is returned per chunk containing all results within that chunk
- **chunk_size** – Size of each chunk to tell InfluxDB to use.
- **dropna** – drop columns where all values are missing
- **data_frame_index** – the list of columns that are used as DataFrame index

Returns: the queried data

Return [ResultSet](#)
type:

```
write_points(dataframe, measurement, tags=None, tag_columns=None,  
field_columns=None, time_precision=None, database=None, retention_policy=None,  
batch_size=None, protocol='line', numeric_precision=None)
```

Write to multiple time series names.

- Parameters**
- :
- **dataframe** – data points in a DataFrame
 - **measurement** – name of measurement
 - **tags** – dictionary of tags, with string key-values
 - **tag_columns** – [Optional, default None] List of data tag names
 - **field_columns** – [Options, default None] List of data field names
 - **time_precision** – [Optional, default None] Either 's', 'ms', 'u' or 'n'.
 - **batch_size** (*int*) – [Optional] Value to write the points in batches instead of all at one time. Useful for when doing data dumps from one database to another or when doing a massive write operation
 - **protocol** – Protocol for writing data. Either 'line' or 'json'.
 - **numeric_precision** – Precision for floating point values. Either None, 'full' or some int, where int is the desired decimal precision. 'full' preserves full precision for int and float datatypes. Defaults to None, which preserves 14-15 significant figures for float and all significant figures for int datatypes.

SeriesHelper

```
class influxdb.SeriesHelper(**kw)
```

Subclass this helper eases writing data points in bulk.

All data points are immutable, ensuring they do not get overwritten. Each subclass can write to its own database. The time series names can also be based on one or more defined fields.

The field “time” can be specified when creating a point, and may be any of the time types supported by the client (i.e. str, datetime, int). If the time is not specified, the current system time (utc) will be used.

Annotated example:

```
class MySeriesHelper(SeriesHelper):
    class Meta:
        # Meta class stores time series helper configuration.
        series_name = 'events.stats.{server_name}'
        # Series name must be a string, curly brackets for dynamic use.
        fields = ['time', 'server_name']
        # Defines all the fields in this time series.
        ### Following attributes are optional. ###
        client = TestSeriesHelper.client
        # Client should be an instance of InfluxDBClient.
        :warning: Only used if autocommit is True.
        bulk_size = 5
        # Defines the number of data points to write simultaneously.
        # Only applicable if autocommit is True.
        autocommit = True
        # If True and no bulk_size, then will set bulk_size to 1.
        retention_policy = 'your_retention_policy'
        # Specify the retention policy for the data points
        time_precision = "h"|"m"|"s"|"ms"|"u"|"ns"
        # Default is ns (nanoseconds)
        # Setting time precision while writing point
        # You should also make sure time is set in the given precision
```

```
classmethod commit(client=None)
```

Commit everything from datapoints via the client.

Parameters: **client** – InfluxDBClient instance for writing points to InfluxDB.

Attention: any provided client will supersede the class client.

Returns: result of client.write_points.

ResultSet

See the [Query response object: ResultSet](#) page for more information.

```
class influxdb.resultset.ResultSet(series, raise_errors=True)
```

A wrapper around a single InfluxDB query result.

error

Error returned by InfluxDB.

```
get_points(measurement=None, tags=None)
```

Return a generator for all the points that match the given filters.

Parameters:

- **measurement** (*str*) – The measurement name

- **tags** (*dict*) – Tags to look for

Returns: Points generator

`items()`

Return the set of items from the ResultSet.

Returns: List of tuples, (key, generator)

`keys()`

Return the list of keys in the ResultSet.

Returns: List of keys. Keys are tuples (series_name, tags)

static `point_from_cols_vals(cols, vals)`

Create a dict from columns and values lists.

Parameters:

- **cols** – List of columns
- **vals** – List of values

Returns: Dict where keys are columns.

`raw`

Raw JSON from InfluxDB.